

Securing AI-Generated APIs at the Speed of AI

Risk-Aware Test Selection and Automated Tests

API WORLD + CLOUDX + AI TECHWORLD 2026

TUESDAY, SEPTEMBER 1, 2026 · 12:00–12:50 PM PDT

WORKSHOP STAGE A (PRO)

Deepankar Das

CEO and Founder, Cyberhead AI

30 min presentation · 15 min live demo · 5 min Q&A



The Core Problem

What AI Agents Can Do

AI agents can now generate API code — routes, schemas, authorization policies, deployment configs — faster than conventional testing and review systems can establish whether that code is safe to ship.

What This Session Covers

This is not a session about whether agents can write an endpoint. They can. It is about what happens next: how a team verifies API correctness, runtime behavior, authorization, tenant isolation, and security at the same pace agents create changes.

📌 This is an engineering and verification session — not a product pitch.

One Change, Six Engineering Lenses

One AI-generated API change crosses all six domains simultaneously.

API Contract

Did we preserve the promise to consumers?

Runtime

Does the application behave correctly under real inputs?

Operations

Can we observe, contain, and recover?

Security

Did we create a new attack or exposure path?

Deployment

Can we safely release and roll back?

Human Accountability

Who accepts the remaining risk?

The Old Testing Rhythm

Built for Human-Paced Change



AI-generated changes arrive continuously. This cadence does not.

Why It Breaks Down

Most API teams inherited a testing rhythm designed for a different delivery model: a major penetration test once or twice a year, broader regression sweeps before a release, and code review in between.

That rhythm was never perfect. It becomes **structurally inadequate** when several agents can create code, tests, configuration, contracts, and deployment artifacts continuously.

- ❑ Penetration testing and regression testing are not unnecessary — they are too late and too infrequent to be the only controls.

The Two Failure Modes

● Run Everything

Full test suite becomes the critical path. Slow feedback, queueing, and weak attribution result. The organization loses the speed it adopted agents to gain.


● Run Too Little

Teams trust unit tests, an agent summary, or a plausible diff. Contract, security, and runtime defects escape: authorization gaps, cache-key leaks, unsafe CORS settings, malformed error paths.

The answer is neither "run nothing" nor "run everything." It is **selective rigor**: the smallest *credible* set of tests for this change, plus non-negotiable invariant checks.

The Workshop Thesis

Test at agent speed by making evidence selective, explainable, and cumulative.



You do not establish trust by slowing agents down until they resemble human developers. The human reviewer receives the full evidence packet and owns the final decision.



The Layered Verification System

Layer	Job	Cadence
Fast per-commit gates	Catch immediate, high-signal defects	Every commit / PR
Impact-selected tests	Cover reachable change surfaces	Every material change
Invariant suites	Protect non-negotiable properties	Every applicable change
Targeted end-to-end checks	Validate cross-layer behavior	High-risk change / pre-merge
Automated penetration tests	Test source-level and runtime security	Change-specific subset + milestone
Milestone validation	Re-establish broad release confidence	Release / major milestone

Each layer has a distinct purpose and cadence — they are not interchangeable test types.

Fast Gates vs. Milestone Validation

Fast Per-Commit Gates

Bounded and fast. Blocks obvious or critical regressions. Runs against the exact diff.

Examples: schema, secrets, targeted tests, critical invariants

Milestone Validation

Broad and slower. Tests integrated and release-level behavior. Re-establishes confidence before release.

Examples: full regression, performance, broader pen-test suite, operational drills

Insisting everything runs for every commit turns every commit into a release candidate. Deferring serious checks until a milestone allows risk to accumulate and makes attribution difficult. The demo will show both layers in action.

Risk-Aware Test Selection

Don't Select from Changed Files Alone

A file-based selector is not enough for APIs. A route can touch a contract. A contract can affect a generated SDK. A query can affect tenant isolation. A cache key can turn a correct query into a cross-tenant data leak.

Formula: Selected corpus = Reachability tests + Invariant floors + Risk-mandated tests

The Reachability Chain

Select from **changed and reachable surfaces**:

- Route → contract → authorization → query → cache
- Service dependencies → generated client
- Config → telemetry

The objective is not to run the fewest tests. It is to run the **smallest credible corpus** for the change.

Invariants Are Non-Negotiable

Risk-aware selection is not permission to optimize security away. Some tests are **invariant floors** — they always run when a relevant surface is touched.



Default-Deny Authorization

If authorization changes, the authorization matrix must run.



Tenant Isolation

If a multi-tenant query or cache changes, isolation tests must run.



Schema & Compatibility

If a shared API contract changes, contract and compatibility checks must run.



Secret Leakage Prevention

Credentials and sensitive values must never appear in source, logs, or artifacts.



A selector may narrow the suite. It may not waive a critical property.



Whitebox Security Validation

Whitebox checks have access to code, configuration, build artifacts, container definitions, and the dependency graph — inspecting patterns a runtime test may not reliably find.

Category	What It Looks For
SQL injection patterns	Unsafe query construction and parameter handling
XSS exposure	Unsafe output or rendering paths
Secret leakage	Credentials in source, logs, or artifacts
Auth coverage gaps	Sensitive routes without expected policy enforcement
Unsafe shell execution	Untrusted input influencing command execution
Container posture	Insecure image, privilege, or runtime settings
Dependency risks	Vulnerable, stale, unpinned, or risky dependencies

Part of a **75-test automated penetration-testing suite**. In the demo, a representative subset is attached to the change based on affected surfaces.

Blackbox API Security Validation

Test the Deployed API as a Client or Attacker Sees It

Source-level correctness does not guarantee deployed behavior. A gateway, proxy, environment variable, CORS policy, or deployment setting can change what the client sees.

Whitebox and blackbox validation are complements. A pass in one is not a waiver for the other.

Category	Example Check
Security headers	Required headers present and correct
Method allowlists	Unsupported methods rejected
Auth enforcement	Missing, invalid, expired credentials fail
Brute-force handling	Repeated failures get defensive handling
Path traversal	Encoded traversal attempts fail
CORS behavior	Only intended origins and methods work
Session cookies	Secure, HttpOnly, SameSite attributes correct

Security Coverage by Cadence

Every Relevant PR

- Fast secret checks
- Auth coverage checks
- Selected security tests
- Targeted blackbox checks

Milestone / Release

- Broader 75-test automated penetration suite
- Deeper dependency and container review
- Performance and operational validation
- Release-path and rollback verification

Security cannot be deferred to an annual event, but the same test depth is not economically sensible for every commit. The system is **continuous** without pretending otherwise.



A Passing Run Is Not Closure

Coverage is not closure.

A dashboard full of green checks does not mean the test system is honest. If a failure occurs and the team reruns until green, the cause may remain unknown. If the same agent that introduced a defect also declares it fixed, the same blind spot may remain.

- 1 Preserve the run artifact
- 2 Investigate the first failure
- 3 Add a counter-example that would catch recurrence
- 4 Re-verify independently when risk warrants it

First Failure and Counter-Examples

When a multi-layer test run fails, preserve the first failure in stable order. Explain it, fix or classify it, rerun the smallest relevant corpus — then add a counter-example.

Assertion	Counter-Example
"Authorization works"	Missing, expired, wrong-role, wrong-tenant credentials
"Cache is safe"	Identical request across two tenants, workspaces, and date ranges
"Contract is compatible"	Existing client fixture against the changed deployment
"Input validation works"	Malformed, oversized, encoded, and boundary inputs

The point of a test is not to confirm the story you already believe. It is to create evidence that could **disprove** it.

Cross-Agent Re-Verification



Agent A
Proposes

Agent B
Challenges

Human
Resolves

Why Correlated Confidence Is a Problem

In agentic workflows, an agent can produce code, tests, and an explanation that all share the same mistaken assumption.

Cross-agent re-verification means a different agent, toolchain, prompt strategy, or model family challenges the proposed change — looking for missing tests, false assumptions, contract inconsistencies, and security gaps.



Independent automated challenge reduces shared blind spots. Human accountability remains.

The Convergence Loop



A failure becomes an **input to the system**, not an interruption to be hidden. We fix the cause rather than only the symptom, rerun the selected corpus, independently challenge the result, and close only when the evidence is durable.

The next fifteen minutes demonstrate this exact loop on a deliberately flawed AI-generated API change.

Live Demo Setup

Sample Application

Multi-tenant usage service

AI-generated change: Add GET

/v1/accounts/{accountId}/usage-summary

Changed surfaces: Route · OpenAPI contract · `usage: read` authorization · tenant query · cache key · SDK · feature flag · telemetry

Seeded defect: Cache key omits account identity

What the Demo Shows

The code will look plausible. Unit tests will mostly be green. But under the right sequence of requests, Tenant B can receive a cached response created for Tenant A.

We will walk through each verification stage, catch the defect, turn it into a durable counter-example, fix it, and produce the evidence packet a reviewer needs.

Demo Stages 1-4: From PR to Invariant Failure

O1

Baseline & Pull Request

View the AI-generated diff: route, schema, scope check, tenant query, cache key, SDK update, telemetry. The flawed cache key is visible but not yet flagged.

O3

Fast Per-Commit Gates

`npm run gate:fast` → PASS in 42s. Schema validates, no secrets, route has auth policy. Useful evidence — but not sufficient.

O2

Change Impact Map

RADAR maps changed surfaces to reachable risks: contract compatibility, authorization default-deny, tenant isolation, cache partitioning, client-generation compatibility.

O4

Impact-Selected Tests + Invariant Failure

RADAR selects 9 tests. 8 pass.

Demo Stages 5–6: Security Checks & Preserving the Failure

Stage 5 — Whitebox & Blackbox Checks

Whitebox PASS: Query parameterization, auth coverage, secret leakage, dependency manifest, no unsafe shell execution.

Blackbox PASS: Security headers, method allowlist, invalid credentials denied, oversized requests rejected, CORS constrained.



Security controls pass while cache partitioning fails. That is exactly why the layered system exists.

Stage 6 — Preserve the First Failure

The run artifact records:

- **First failure:** `cache.usage_summary.key_partition`
- **Violated invariant:** `tenant_isolation`
- **Observed behavior:** Tenant B received a cached response populated by Tenant A
- **Suspected cause:** Cache key lacks `accountId`
- **Required counter-example:** Same workspace and date range across distinct accounts
- **Closure state:** open

Demo Stage 7: Fix, Counter-Example & Re-Verification

The Fix

The corrected cache key adds `accountId` and `role` to the boundary:

```
usage-summary:${accountId}:${workspaceId}:${role}:  
${startDate}:${endDate}
```

The Counter-Example Test

The new test verifies that Tenant A and Tenant B each receive their own data for identical query parameters — permanently encoding the class of failure, not just the exact request.

Independent Re-Verification

```
npm run radar:run → 9/9 PASS
```

```
npm run verify:independent challenges:
```

- Route and cache-key boundaries
- Account, workspace, role, and date isolation
- Counter-example presence confirmed
- No additional contract or auth coverage gap found



Independent verdict: PASS



Demo Stage 8: Final Receipt & Release Decision

The final run artifact gives the reviewer a complete, reviewable evidence packet:

What Changed

8 surfaces: route, contract, auth policy, tenant query, cache key, SDK, feature flag, telemetry

What Was Tested

9 selected tests · 5 invariant floors · whitebox PASS · blackbox PASS · independent verification PASS

Failure Resolved

Cache partition failure closed · counter-example added · first-failure artifact preserved

Release Decision

Admissible for merge. Milestone still requires: broader pen suite, load test, feature-flag rollout, rollback verification



Admissible for merge ≠ ready for unrestricted rollout. Human reviewer owns the final decision.

The System in One Picture

```
graph TD; A[AI-generated change] --> B[Fast gates]; B --> C[Impact map + invariant floors]; C --> D[Selected correctness and security tests]; D --> E[Run artifact + first-failure discipline]; E --> F[Counter-example + cross-agent re-verification]; F --> G[Human approval]; G --> H[Milestone validation and staged release];
```

AI-generated change
↓
Fast gates
↓
Impact map + invariant floors
↓
Selected correctness and security tests
↓
Run artifact + first-failure discipline
↓
Counter-example + cross-agent re-verification
↓
Human approval
↓
Milestone validation and staged release

The Operating Model

Not annual penetration testing. Not a full suite on every diff. Not trust in an agent's summary.

Layered validation

Each layer has a purpose and cadence.

Risk-aware selection

Smallest credible corpus for the change.

Durable failure closure

Counter-examples, not reruns.

Human accountability

Evidence packet, not confidence theater.

The goal is not to remove human judgment — it is to reduce the amount of **unstructured, unsupported** judgment a reviewer must make.

Q&A — Anticipated Questions

Why not run the full suite on every change?

Use the full suite where broad assurance is justified — especially milestones. At agent speed, per-change verification also needs fast, attributable, risk-aware feedback. A full suite is often too slow and too weakly tied to the precise change.

How do you know the selected corpus is sufficient?

Combine reachability with mandatory invariant floors, risk-mandated tests, milestone validation, and a learning loop that turns escaped failures into new tests or rules.

Does this replace penetration testing?

No. It makes security validation more continuous and change-aware. Broader penetration testing and human security review remain important milestone controls.

What should never be automated?

Risk acceptance, business authorization, customer-impact trade-offs, ambiguous incident decisions, and final approval of high-blast-radius or irreversible production actions.

Questions

What sensitive API surface would you use as the first pilot?

Start with one repository, one invariant suite, one run artifact, and one class of AI-generated change.

Start Small

Pick one sensitive API module and define its invariants.

Preserve Evidence

Require run artifacts and counter-examples for every failure.

Build the Loop

Use a simple change-to-test selection record in pull requests and grow from there.

