

Deepankar Das

[My Events](#)

[API World + CloudX + AI TechWorld 2026](#)

Securing AI-Generated APIs at the Speed of AI: Risk-Aware Test Selection and Automated Tests

API World + CloudX + AI TechWorld 2026 in 6 days

Congratulations! Your **session has been accepted!**

You have confirmed your participation on 15 Jun 2026.

Securing AI-Generated APIs at the Speed of AI: Risk-Aware Test Selection and Automated Tests

Tuesday, 1 Sep 2026 12:00 pm - 12:50 pm (50 minutes)

in Pacific Daylight Time (UTC-07:00)

API World -- Workshop Stage A (PRO)

Description

Annual penetration tests and occasional regression sweeps are not enough when AI-generated code lands continuously. As part of a multi-agent engineering workflow, we needed API validation that could keep pace with code generation while still catching real security, contract, and runtime failures before they reached production.

We built a layered verification system for a production multi-tenant API platform: fast per-commit gates, impact-selected tests, invariant suites for sensitive modules, targeted end-to-end checks, and a 75-test automated penetration testing suite spanning whitebox and blackbox validation. The whitebox side checks for

SQL injection patterns, XSS exposure, secret leakage, auth coverage gaps, unsafe shell execution, weak container posture, and dependency risks. The blackbox side validates security headers, method allowlists, oversized payload behavior, auth enforcement, brute-force handling, path traversal resistance, CORS behavior, session cookie attributes, and other API-facing controls.

The deeper lesson is not just coverage, but closure discipline. We will show how we used run-artifact metrics, first-failure reruns, counter-example checks, and cross-agent re-verification to keep the test system honest when AI-generated changes broke multiple layers at once. The talk also covers what we did not automate and why, so teams can focus human effort where automation still falls short.

Attendees will leave with a practical model for testing AI-generated APIs at production speed, a concrete way to divide fast gates from milestone validation, and reusable patterns for combining API correctness, runtime safety, and security testing into one coherent engineering system.

event Tracks	API Design, Development, and Deployment	API Operations and Management	API Security	APIs in the Age of AI
--------------	---	-------------------------------	--------------	-----------------------

Additional Topics	API Management	Low code/No code	Model Context Protocol (MCP)	Agentic AI
-------------------	----------------	------------------	------------------------------	------------

Session format	Technical Workshop/Tutorial (50 Min)	Technical Session (25 Min)	Best Practices Session (25 Min)
	Business/Thought Leadership (25 Min)		

Submitted

16 May 2026 3:10 pm



SESSION OWNER

Deepankar Das

CEO and Founder @ CyberHead AI

Deepankar Das is a technology executive and agentic AI leader with over two decades of experience building and scaling enterprise software, cloud platforms, and AI-driven products for global customers. He has led engineering organizations of 50–400+ people across companies like Dell and Persistent Systems, and now serves as CEO and Founder of CyberHead AI, where he is building agentic AI platforms for the industry verticals. Deepankar's work focuses on practical, production-grade Agentic Engineering: designing multi-agent systems, safety guardrails, and orchestration patterns that can be trusted in regulated, high-stakes environments. He has architected AI and agentic workflows that integrate deeply with existing enterprise stacks—cloud-native infrastructure, distributed data platforms, and complex line-of-business applications—while meeting stringent requirements around security, compliance, and reliability. Throughout his career, Deepankar has taken multiple products from 0→1 and then scaled them to large, global customer bases, combining strong technical judgment with hands-on leadership. His recent work includes building AI coding assistants and guardrail layers that safely accelerate engineering teams, as well as RAG-based and agentic systems that augment knowledge workers in domains such as legal, finance, and infrastructure operations. Deepankar brings a pragmatic, operator's perspective to Agentic AI: what it really takes to move from impressive demos to robust systems that teams can depend on every day. He speaks frequently about agentic architectures, multi-tenant SaaS design, secure data and compliance patterns for AI, and how to evolve traditional engineering organizations into AI-native ones without losing rigor or control.